

Simplified Neutreeko using Reinforced Learning

André Gomes - 201806224
Gonçalo Teixeira - 201806562
Luís Recharte - 201806743

Specification of Reinforced Learning Problem

Game Logic

- Create Starting Condition
- Possible moves
- Manage player turns
- Perform a move and update the board
- Verify end game and tie conditions

Environment

- Step - Apply the Agent action to the environment
- Reset - Reset the env
- Render - Print a representation of the env
- Close - Finish the episode
- Done - Check if episode is done

Agents

- **Random choice** Agent
- **Q-learning** Agent
- **SARSA Agent**
- Monte Carlo Agent (not implemented)

Related Work & References

Board games with OpenAI Gym:

- Abalone - <https://github.com/towzeur/gym-abalone>
- Go - <https://github.com/aigagror/GymGo>

Spaces' definition -

<https://github.com/openai/gym/tree/master/gym/spaces>

Create an environment -

<https://github.com/openai/gym/blob/master/docs/creating-environments.md>


Table of environments -

<https://github.com/openai/gym/wiki/Table-of-environments>

States, Observation and Action Spaces in Reinforcement

Learning - <https://medium.com/swlh/states-observation-and-action-spaces-in-reinforcement-learning-569a30a8d2a1>

Q-Learning Agent - <https://medium.com/swlh/introduction-to-q-learning-with-openai-gym-2d794da10f3d>



abalone +





Tools and Algorithms

Anaconda environment -
Python 3.8

JetBrains IntelliJ

Gym, Numpy, Matplotlib

Q-Learning and SARSA

Work carried out

1

Started the structure of a OpenAI Gym project

2

Implemented the game logic

3

Implementation of the game loop and the environment. Random Agent to play the game

4

Define the rewards and their values. Implement the remaining agents

5

Transpose code to a Jupyter notebook



Sinalizar para seguimento.



Jan Kristian Haugland <admin@neutreeko.net>

sáb, 22/05/2021 12:37

Para: André Gomes



Hi,

I am not entirely sure how I arrived at 3,450,515 (this would date back around 19 years). I have probably assumed that the Next player (N) does not already have three in a row, which gives them $C(25, 3) - 48 = 2252$ possible positions for their pieces, where $C(n, k)$ denotes the binomial coefficient (n over k). The Previous player (P) should then have $C(25 - 3, 3) = 1540$ possible position for each of them, for a total of $2252 \times 1540 = 3,468,080$ positions. Most likely I have subtracted positions that are impossible to reach because there is no position from which the Previous player could have made a legal move and reached the current position. Two examples are

```
00000
0P0P0
00000
NN0P0
ON000
```

and

```
00000
ONPPN
00000
N00P0
00000
```

But I guess for this purpose, this is not extremely important, and we can use all the 3,468,080 positions.

I am not sure what would be a useful way to map each position to a unique number. Perhaps one could maintain a list of the $C(25, 3) = 2300$ ways to place three pieces of one colour, for a total of $2300^2 = 5,290,000$ positions with both colours, and then sift out those that collide, if that is feasible. This is what my program does. By the way, on my web site


<https://www.neutreeko.net/neutreeko.htm>

Rewards:
Win: 20
Step: -1

Random
Agent



Q-Learning Agent



SARSA Agent

